Assembling Close Strains in Metagenome Assemblies Using Discrete Optimization

Tam Khac Minh Truong[1,*], Roland Faure[1,2,*] [a] and Rumen Andonov[1] [b]

[1]*Univ. Rennes, INRIA RBA, CNRS UMR 6074, Rennes, France*

[2]*Service Evolution Biologique et Ecologie, Université Libre de Bruxelles (ULB), 1050 Brussels, Belgium*

Abstract: Metagenomic assembly is essential for understanding microbial communities but faces challenges in distinguishing conspecific bacterial strains. This is especially true when dealing with low-accuracy sequencing reads such as PacBio CLR and Oxford Nanopore. While these technologies provide unequaled throughput and read length, the high error rate makes it difficult to distinguish close bacterial strains. Consequently, current de novo metagenome assembly methods excel to assemble dominant species but struggle to reconstruct low-abundance strains. In our study, we innovate by approaching strain separation as an Integer Linear Programming (ILP) problem. We introduce a strain-separation module, strainMiner, and integrate it into an established pipeline to create strain-separated assemblies from sequencing data. Across simulated and real experiments encompassing a wide range of error rates (5-12%), our tool consistently compared favorably to the state-of-the-art in terms of assembly quality and strain reconstruction. Moreover, strainMiner substantially cuts down the computational burden of strain-level assembly compared to published software by leveraging the powerful Gurobi solver. We think the new methodological ideas presented in this paper will help democratizing strain-separated assembly.

1 INTRODUCTION

In the field of metagenomic sequencing, the analysis of bacterial communities is a complex undertaking, complicated by the presence of conspecific strains (i.e. strains of the same species). Current *de novo* metagenome assembly methods can reconstruct the chromosomal sequences of prevalent species but generally struggle to produce strain-level reconstructions. This capability is vital for discerning subtle genetic differences among microorganisms that hold crucial functional significance in different environments. For instance, many *Escherichia coli* strains are commensal, while others are pathogenic (Frank et al., 2011).

The challenge posed by the "strain separation problem," as outlined in (Vicedomini et al., 2021), arises from two primary factors: the uncertain and potentially substantial quantity of conspecific strains and their uneven distribution within the sample. In this study, we will use the term "haplotype" to de-

scribe each strain's genome. It is important to note that this problem isn't perfectly defined because the concept of a "strain" lacks absolute clarity, given that any two individuals naturally exhibit some genetic distinctions. Within the scope of this research, we consider a haplotype to be a contiguous sequence of nucleotides present in sufficient abundance within the sample. This approach aligns with similar practices found in other relevant works, including (Vicedomini et al., 2021).

Numerous studies have focused on addressing the strain separation problem primarily in the context of short-read sequencing. Methods like DESMAN (Quince et al., 2017), STRONG (Quince et al., 2020), ConStrains, LSA (Cleary et al., 2015), OPERA-MS (Bertrand et al., 2019), SAVAGE (Baaijens et al., 2017), and strainXpress (Kang et al., 2022) have made notable attempts to tackle this problem, albeit with various limitations. Additionally, StrainPhlAn (Truong et al., 2017) and StrainEst (Albanese and Donati, 2017) have introduced references-based techniques to profile communities and finely distinguish known strains.

However, with the advent of long-read metage-

[a] https://orcid.org/0000-0003-2245-4284

[b] https://orcid.org/0000-0003-4842-7102

*These authors contributed equally to this work.

nomic sequencing, previous methods designed for short-read data are no longer suitable due to fundamental differences in data characteristics. Existing long-read assemblers, including Canu (Koren et al., 2017) and metaFlye (Kolmogorov et al., 2020), have emerged as state-of-the-art solutions for metagenome assembly but are not designed to fully address the strain separation challenge.

PacBio High-Fidelity (HiFi) sequencing technology has emerged as a promising solution for metagenomic strain separation, as the extremely low error rate can help distinguish very similar sequences. Specialized software such as hifiasm-meta (Feng et al., 2022), strainFlye (Fedarko et al., 2022) and stRainy (Kazantseva et al., 2023) have been developed to exploit this kind of data. However, achieving sufficient coverage to recover rare strains in complex metagenomes is challenging and expensive. In this article we will focus on higher-error rates technologies such as Oxford Nanopore and Pacbio CLR, which are more affordable and have higher throughput.

MagPhase (Bickhart et al., 2022) and iGDA (Feng et al., 2021) have been proposed to phase Single Nucleotide Polymorphisms (SNPs) in metagenomes. However, these tools were not meant to recover full haplotypes, but rather phase specific regions of genomes. MagPhase only returns lists of SNPs, neglecting more complex variants, while iGDA failed to run on the full metagenomes on which we tested it.

Strainberry (Vicedomini et al., 2021) has been the first tool to propose a solution to tackle the strain separation problem at the scale of full metagenome assembly. It applies iteratively HapCUT2 (Bansal, 2022), a tool developed for diploid phasing. The authors proved that Strainberry was capable of phasing simple genomes, but the tool is intrinsically limited to simple metagenomes, i.e. no more two or three conspecific strains. More recently, the software Hair-Splitter (Faure et al., 2023) has been introduced. Hair-Splitter begins by aligning all the reads on a first draft assembly and calls polymorphisms. It then clusters reads by similarity, aiming to obtain one cluster per strain. Finally, it re-assembles the reads to obtain the phased assembly.

In this article, we introduce a new method to separate strains in metagenome assemblies using error-prone long reads. This method introduces a completely new methodological tool by framing strain separation as an Integer Linear Programming (ILP) problem, allowing us to harness the power of the Gurobi solver (through gurobipy v.10.0.3, gurobi.com) to achieve high computational efficiency (Gurobi Optimization, LLC, 2023). We implement this method in an algorithm named strainMiner. We


Figure 1: The strainMiner pipeline: (i) Reads are aligned on the reference or draft assembly, (ii) on each window of the assembly, reads are separated by haplotype of origin - only three windows are shown here, (iii) all groups of reads are locally reassembled and (iv) the locally reassembled contigs are scaffolded to produce longer contigs.

integrate it in the HairSplitter pipeline to obtain a complete software capable of producing a new assembly, which we will call "the strainMiner pipeline". We show that on mock communities (based on real and simulated data) the strainMiner pipeline improves over Strainberry and compares favorably with Hair-Splitter to phase metagenome assemblies in terms of assembly completeness, while being an order of magnitude faster (than Strainberry) or more memory-efficient (than HairSplitter).

2 CONTRIBUTION

2.1 Pipeline

The strainMiner pipelines takes as input a draft assembly or a reference genome and a set of long reads and aims at producing a strain-separated assembly. It comprises four primary stages and is illustrated Figure 1. These stages include: (i) aligning the reads to the draft assembly or reference genome, (ii) the strainMiner algorithm - clustering the reads based on their respective haplotypes, (iii) conducting haplotype-specific assembly of the reads, and (iv) enhancing assembly contiguity through scaffolding. Excluding step 2, this pipeline mirrors the one detailed in (Faure et al., 2023).

Our contribution lies in an original method to propose a solution for the second step of the pipeline, i.e. separating aligned reads by haplotype of origin. The rest of the pipeline was forked from (Faure et al., 2023), replacing a native HairSplitter module with a module of our own.

The specific challenge we are addressing can be described as follows: given a set of reads aligned to a reference sequence, distinguish groups of reads according to their haplotype of origin. In an ideal sce-

nario, all the reads of the dataset originating from the same strain would be grouped together. However, achieving this level of separation across an entire genome is not always feasible. Indeed, it is impossible to phase two consecutive variants when they are too far apart to be spanned by at least one read. Therefore, our task focuses on separating reads *locally* by their haplotype of origin. Here we choose to cut the genome in windows of length w, where w should be smaller than the average read length (5000 by default).

2.2 Intuition

The intuition behind strainMiner is similar to the one behind (Faure et al., 2023), originally introduced in (Feng et al., 2021).

The crucial information for segregating reads by their haplotype of origin is found within the polymorphic sites—specific genomic locations where different strains exhibit variations. However, it is impossible to distinguish on a given read a sequencing error from a variation at a polymorphic locus. Using the pileup of read might give us an indication, if many reads show the same nucleotide at the same locus. Nevertheless, this approach has its limitations, as alignment artifacts can also create such loci in the pileup, and the signal from rare strains may not be very pronounced. The trick lies in examining multiple loci simultaneously, as reads originating from a particular strain will consistently display similar behavior at polymorphic sites. Conversely, the base pileup between two non-polymorphic sites should not exhibit significant correlation since alignment or sequencing errors occur randomly among all the reads.

As a result, strainMiner endeavors to identify a set of positions which share identical patterns of variation. We expect these to represent, for example, all loci bearing variants specific to one strain.

This strategy allows us to detect even rare strains effectively. For instance, consider a hypothetical scenario involving a mixture of two strains, where strain A constitutes 99% of the mix and strain B a mere 1%. Consider also a collection of a thousand reads spanning two polymorphic sites, denoted as a and b. In an ideal, error-free pileup, conducting a chi-square test for independence with one degree of freedom between the two loci yields a p-value smaller than 10^{-215}.

However, the introduction of errors significantly diminishes the statistical power. In our simulations, we introduced random substitution errors with a probability of $p = 0.1$ for all bases across 10,000 simulations. The p-value for the correlation between a

and b remained low, averaging 10^{-16} and reaching 10^{-6} in the most unfavorable scenario. Nevertheless, it is important to consider that thousands of non-polymorphic positions can potentially correlate with a in one pileup. We need also to emphasize that alignment artifacts can introduce more complex errors with locally higher error rates, further reducing the statistical power of correlation.

The solution is to include more loci to drastically reduce the risk of spurious correlations. In this same example, we introduced a third locus, denoted as c, while maintaining the 0.1 error rate. We applied the one-degree of freedom chi-square test to assess the relationship between the three positions. This time, the probability of encountering three non-polymorphic positions with correlations as strong as those observed between a, b, and c by chance was found to be below 10^{-200} in all 10,000 simulations. While this example simplifies the complexities of pileup errors, it emphasizes two fundamental aspects of the method: a) the joint observation of multiple loci significantly enhances the statistical power to distinguish between errors and polymorphism, and b) even low-abundance strains can be reliably identified.

2.3 Preprocessing

In each window, strainMiner considers only reads that span at least 60% of the window's length. Subsequently, strainMiner transforms the read pileup into a binary matrix, where each row corresponds to a read, each column corresponds to a position, and cell (i, j) contains the number one if the base at position j of read i matches the dominant base at that position, the number zero if it matches the most common alternative allele, and remains empty otherwise. Empty cells can occur if a read doesn't cover a position or if the base in a read at a given position isn't among the two most common bases at that position.

Next, columns are filtered to retain only those in which the most common base is present in less than a proportion p of the aligned reads at that position. By default, strainMiner sets p to 0.95, striking a balance between computational efficiency and precision. However, users seeking to ensure the recovery of low-abundance strains can set p to a higher value.

To populate the empty cells, strainMiner implements the well-known K-nearest-neighbor imputation strategy (Fix and Hodges, 1989), used widely, for example in (Troyanskaya et al., 2001). It identifies for each read its "nearest neighbors," i.e. reads with the smallest hamming-distance proportionally. Then, for each empty cell in a row, strainMiner uses a majority vote from the five closest neighbors that have

non-empty cells at that position to impute the missing value.

The result is a binary matrix filled with zeroes and ones. Our primary objective is to identify groups of positions where reads exhibit similar behavior, hence strainMiner essentially finds sets of similar columns in the matrix. Each of these groups will represent a bipartition of strains, distinguishing strains with the predominant allele from those with the secondary allele.

In practice, this task is considerably more challenging than it may initially appear. Attempting to identify position groups through a straightforward hierarchical clustering approach often proves ineffective. The reason behind this lies in the calculation of pairwise distances between two positions, which can be deceptive. In many cases, sequencing and alignment errors can overshadow small differences resulting from low-abundance strains. As emphasized earlier, the key is to jointly examine more than just two loci simultaneously.

2.4 Finding Bipartitions of Reads

The approach used to identify bipartitions is depicted in Figure 2. It begins by identifying the largest bicluster (a cluster of row and columns) of ones in the matrix, allowing a small number of zeros to be tolerated to account for sequencing errors (referred to as a quasi-bicluster of ones). Figure 2.b visualises the largest quasi-bicluster found in the first step. Next, another submatrix is formed, including the columns from the first quasi-bicluster and all the reads not included in the initial one. In this new submatrix, the largest quasi-bicluster of zeros is sought (Figure 2.c). This process continues, alternating between quasi-biclusters of ones and zeros, until all reads have been included in these clusters or all columns excluded from the biclusters. The columns shared by all the clusters define a bipartition of reads (Figure 2.e). These columns are excluded from the total matrix and the process is repeated to find another bipartition of reads using the remaining columns.

Once a list of bipartitions has been determined, reads are split in groups following a very simple rationale: two reads are in the same group if and only if they are grouped together in all bipartitions.

2.4.1 Finding Largest Submatrix

In strainMiner, the primary problem is to discover the largest bicluster during each iteration. We accomplish this by employing a Integer Linear Problem (ILP) model. This model's purpose is to select a group of rows and columns where you can find the

highest count of a specific value (either 0s or 1s). This problem is known as finding a maximum edge biclique problem and it has been proven to be NP-complete (Peeters, 2003). In strainMiner, we search for a maximum edge quasi-biclique that tolerates a small amount of errors.

In the context of a matrix $A \in Z_2^{|U| \times |V|}$ with coefficients being 0 or 1, and with a set of rows U and a set of columns V, we use binary variables x_{ij}, u_i and v_j to denote cell selection, row selection, and column selection, respectively. Here, a binary variable equals 1 to indicate the selection of a cell, row, or column, and 0 otherwise.

The following Integer Linear Program is then formulated:

$$\max \sum_{i \in U} \sum_{j \in V} A_{i,j} x_{ij}, \tag{1}$$

$$x_{ij} \leq u_i, \ \forall i \in U, \forall j \in V \tag{2}$$

$$x_{ij} \leq v_j, \ \forall i \in U, \forall j \in V \tag{3}$$

$$x_{ij} \geq u_i + v_j - 1, \ \forall i \in U, \forall j \in V \tag{4}$$

$$\sum_{i \in U} \sum_{j \in V} (1 - A_{i,j}) x_{ij} \leq \varepsilon \times \sum_{i \in U} \sum_{j \in V} x_{ij} \tag{5}$$

$$u_i, v_j \in \{0, 1\}, \ x_{ij} \in \{0, 1\} \ \forall i \in U, \ \forall j \in V \tag{6}$$

The function to maximize, (1), counts for the number of ones in a submatrix determined by the binary variables having value 1. Constraints (2), (3), (4) mean that cell A_{ij} is selected into the solution (i.e. $x_{ij} = 1$) if and only if its corresponding row i and column j are also chosen into the solution (i.e. $u_i = 1$ and $v_j = 1$).

The coefficient A_{ij} represents the value of the cell at position i and j, when searching for occurrences of 1s in the matrix, A_{ij} is used directly. However, if the search is for 0s, the coefficient is reversed to $(1 - A_{ij})$:

$$\max \sum_{i \in U} \sum_{j \in V} (1 - A_{i,j}) x_{ij} \tag{7}$$

Constraint (5) ensures that the submatrix contains at least a proportion $1 - \varepsilon$ of ones. This constraint can be reversed to ensure a minimum proportion of zeros:

$$\sum_{i \in U} \sum_{j \in V} A_{i,j} x_{ij} \leq \varepsilon \times \sum_{i \in U} \sum_{j \in V} x_{ij} \tag{8}$$

To identify a bipartition, strainMiner involves solving a series of this ILP models using the Gurobi solver, alternating between looking for ones and zeros. This process is repeated to discover new bipartitions until the model identifies a submatrix with less than 5 rows or 5 columns — a threshold empirically set to identify a statistically significant signal. When this condition is met, the algorithm terminates and returns the identified bipartitions.
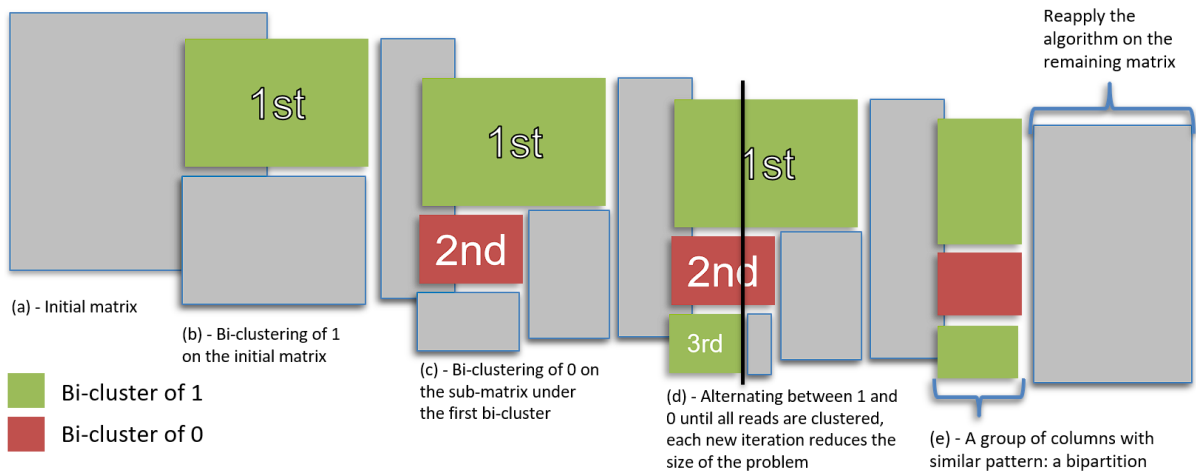
Figure 2: Strategy implemented to find bipartitions.

2.4.2 Acceleration with Divide and Conquer

In order to reduce the computational cost as well as the runtime of strainMiner, we implement a strategy aiming to reduce the size of the matrices prior to running the model.

If a straightforward column-wise hierarchical clustering cannot be used to make fine distinctions between bipartitions (for reasons explained above), it remains an effective strategy for distinguishing highly divergent positions. Thus, a basic hierarchical clustering of the positions with Hamming distance and complete linkage is performed. The clusters are formed by splitting vertically the initial matrix, a threshold of 35% is set as the distance threshold. Clusters with the linkage distance higher than the limit will not be merged. This means that two columns in two different clusters have more than 35% divergence and will never be even considered to be grouped in the same bipartition - we purposefully chose a conservative threshold.

The distance threshold needs to balance creating separate sub-matrices while preserving smaller patterns. If the threshold is too low, even minor data variations and errors can wrongly divide related positions, breaking patterns. If it is too high, sub-matrices that could be split get grouped, forming fewer, larger sub-matrices and the computational gain is not optimal.

The resulting sub-matrices are subsequently assessed for "ambiguity". A submatrix consisting of full rows of zeros or ones is by itself a bipartition c.f. Figure 3 right. The other submatrices are "ambiguous" c.f. Figure 3 left. Bipartitions are sought only in the ambiguous matrices, using the algorithm described above.
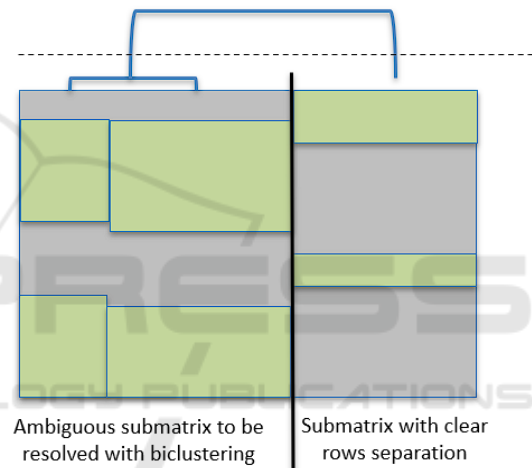

Figure 3: Hierarchical clustering is performed to cut the matrix into smaller distinct sub-matrices defining distinct bipartitions.

3 RESULTS

3.1 Datasets

We benchmarked strainMiner on three datasets of increasing complexity. The first one is a mix of five *Vagococcus fluvialis* strains sequenced in (Rodriguez Jimenez et al., 2022). The bacteria were sequenced using a R9.4.1 Nanopore flowcell. The reads were barcoded by strain of origin, allowing us to assemble the five strains separately. By ignoring the barcodes, we obtain our dataset of five mixed strains of approximately equal abundance sequenced with Nanopore. Three of these strains were almost identical, approximately turning the problem in a three-strain separation problem with one predom-

Figure 4: Comparison of Strainberry, HairSplitter and strainMiner on the three real sequencing experiments in terms of, from left to right, completeness of assembly, CPU time and maximum memory usage. The legend applies to the three plots.

inant species.

The second dataset is the sequencing of the Zymobiomics gut microbiome standard by Nanopore R9.4.1 (Q9 reads) and Nanopore 10.4.1 (Q20+ reads) flowcells (accession numbers SRR17913199 and SRR17913200). The mix contains a total of 21 strains of bacteria, archea and yeast. In particular, it contains five *Escherichia coli* strains in equal abundance, for which we will compare strain separation techniques. This is the dataset having a known solution with highest number of conspecific strains we could find.

In a final step, we conducted simulations to gauge the limits of strainMiner's strain separation capabilities. We adapted the simulation protocol outlined in (Vicedomini et al., 2021) and (Faure et al., 2023) to investigate how the number of strains and the depth of coverage influence the effectiveness of strain separation. We downloaded reference genomes for ten distinct strains of *Escherichia coli* from the SRA and proceeded to simulate Nanopore reads with a 5% error rate and 50x coverage using the default "Nanopore2023" mode of Badreads (Wick, 2019). These reads were then combined to emulate metagenomic sequencing scenarios. The ten strains mirrored those used in (Faure et al., 2023) and included 12009 (GCA_000010745.1), IAI1 (GCA_000026265.1), F11 (GCA_018734065.1), S88 (GCA_000026285.2), Sakai (GCA_003028755.1), SE15 (GCA_000010485.1), *Shigella flexneri* (GCF_000006925.2), UMN026 (GCA_000026325.2), HS (GCA_000017765.1), and K12 (GCF_009832885.1). The simulated reads are available at https://zenodo.org/records/10362565. To assess the impact of the number of strains on assembly completeness, we assembled and separated mixtures of 2, 4, 6, 8, and all 10 strains. Additionally, to evaluate the influence of coverage on assembly completeness, we downsampled the 12009 strain to 30x, 20x, 10x, and 5x within the context of the 10-strain mixture, measuring the 27-mer completeness

of the 12009 strain in the various assemblies.

All of these datasets were assembled using metaFlye (Kolmogorov et al., 2020) because it is to our knowledge the only noisy long-read assembler specialized in metagenome assembly. The metaFlye assembly was then run through the three software.

3.2 Evaluation Metrics

To evaluate the quality of obtained assemblies we measured assembly length, N50, misassemblies, mismatches and indels, measured by software metaQuast (Mikheenko et al., 2015). We used options –unique-mapping and –reuse-combined-alignments in metaQuast to avoid a sequence (a contig, or part of it) to be mapped on multiple distinct reference locations. The evaluation of metagenome assemblies in presence of highly similar references is however a challenging task. As metaQuast is based on sequence alignment, it could suffer from a sub-optimal mapping of contigs to the references. For this reason, we decided to complement metaQuast's metrics by computing the k-mer completeness (k=27) with KAT (Mapleson et al., 2016).

3.3 Benchmark

All results were obtained by running the software on a server housing 16 Intel Xeon CPUs with four cores each, running at 2.7GHz. 3.1 TB of RAM was available.

3.3.1 Assembly Evaluation

The summarized metaQuast metrics in Table 1 show that the strainMiner pipeline yields highly precise assemblies, exhibiting fewer misassemblies, mismatches, and indels compared to Strainberry or HairSplitter. However, this enhanced precision is balanced by lower contiguity. To manage the size of Table 1, we haven't included all the *E. coli* experiments,

Figure 5: 27-mer completeness of the 12009 strain in a 10-strain mix. Coverage of the 12009 strain on *x* axis, the nine other strains are covered at 50x.

but the downsampling experiments displayed similar statistics to the 10-strains mix.

On the Zymobiomics and the *Vagococcus fluvialis* datasets, the strainMiner pipeline proved the best software to recover strain diversity as measured by 27-mer completeness, though the performances of Hair-Splitter were almost equivalent (Figure 4). Notably, the strainMiner pipeline yielded assemblies of equivalent quality when using Q9 and Q20+ chemistry, underscoring its robustness in managing variations in read quality.

In our experiments with *E. coli* (Figure 6), strain-Miner consistently demonstrated its capability to separate a substantial number of strains, maintaining a stable 27-mer completeness across scenarios involving 2, 4, 6, 8, and 10 strains. The strainMiner pipeline consistently outperformed HairSplitter to a slight extent in all mix scenarios.

The coverage reduction experiment depicted in Figure 5 revealed that both strainMiner and Hair-Splitter exhibited an impressive capability to retrieve strains with low coverage. Even when the strain's coverage dropped to as low as 5x, representing only 1.1% of the total mix, these tools managed to recover a substantial portion of the previously lost 27-mers. There remained a noticeable positive correlation between the coverage level of the downsampled 12009 strain and its resulting completeness.

The performance of Strainberry is difficult to assess, as it performed extremely well in 2-strain, 4-strain, 8-strain and 20x-downsampling mixtures but failed to produce good strain-separated assemblies in the other *E. coli* scenarios. This inconsistent behavior is probably due to it being conceived for relatively simple metagenomes (i.e., up to five strains).

These experiments offer insights indicating that all three software solutions can attain assemblies of comparable quality, albeit with varying tradeoffs between contiguity and accuracy, when employing high-quality Nanopore reads. Figure 4 and additional ex-

aminations performed on mixtures of *E. coli* with more error-prone data (not shown) suggest that Strainberry encounter difficulties when confronted with higher error rates.

3.3.2 Resource Usage

All the times reported are total CPU times. Stain-Miner is trivially parallel, like HairSplitter and Strainberry, but in practice we ran it on only one thread because we were limited by the Gurobi license.

Across all datasets, strainMiner consistently exhibited a processing speed more than tenfold faster than Strainberry, while its runtime was approximately on par with that of HairSplitter (Figures 4, 6). For all three approaches we expect a linear increase in runtime when increasing the length of the assembly and the number of species, as aligning is an almost-linear process in the length of the reference and contigs are processed independently by the strain-separation software.

On the real sequencing data, strainMiner used between 7 and 30 times less peak memory than Hair-Splitter (Figure 4).

As a whole, strainMiner significantly diminishes the memory usage of the HairSplitter pipeline without impacting negatively on its speed.

4 DISCUSSION

In this manuscript, we introduce strainMiner, an innovative approach aimed at enhancing the performance of the read-separation module within the HairSplitter strain-separation pipeline, leading to strain-aware metagenome assemblies. For the first time, we frame the strain separation as an ILP problem. This allows us to use a well-established, highly optimized solver to tackle the intensive computations needed for such a task. By helping the solver with simple preprocessing techniques inspired from data mining, the final pipeline is considerably more frugal in time and memory than the state-of-the-art software. This development presents new opportunities for biologists, as strain separation is a computationally demanding task, and having access to such computational resources can be a significant constraint for biology laboratories. On real datasets, the strainMiner pipeline also compared favorably to published software in terms of strain recovery and arguably provided the most accurate assemblies. Additionally, we showed that strainMiner could perform well on long reads with varying error rate without any manual fine-tuning.

Figure 6: Comparison of Strainberry, HairSplitter and strainMiner on mixes of an increasing number of *E. coli strains* in terms of, from left to right, completeness of assembly, CPU time and maximum memory usage. Note that CPU time is represented in logarithmic scale. The legend applies to the three plots.

Table 1: metaQuast metrics measuring the quality of assemblies on the three real and two simulated datasets. The best value or best values among of the three strain-separated assemblies (excluding the original assembly) is in bold font.

		Total length (Mb)	N50 (kb)	Misassemblies	Mismatches /100kbp	Indels /100kbp
Vagococcus	metaFlye	4.35	162	40	340.14	383.57
(14 Mb)	Strainberry	5.51	**135**	135	77.71	510.25
	HairSplitter	**9.43**	82	117	89.93	**327.91**
	strainMiner	8.63	30	**48**	**50.8**	340.26
Zymo Q9	metaFlye	61.9	1904	110	91.14	56.28
(78 Mb)	Strainberry	62.8	**225**	122	133.04	109.85
	HairSplitter	**86.9**	74	170	124.54	124.54
	strainMiner	**68.4**	**225**	**95**	**69.8**	**55.44**
Zymo Q20	metaFlye	49.8	383	126	109.51	75.79
(78 Mb)	Strainberry	59.7	**121**	141	109.12	67.23
	HairSplitter	**66.2**	84	142	75.42	**61.23**
	strainMiner	59.8	51	**109**	**67.97**	65.14
E. coli	metaFlye	11.6	56	79	469.27	277.71
10 strains	Strainberry	18.3	**63**	**175**	363.06	137.01
(48 Mb)	HairSplitter	50.5	48	812	266.55	**68.04**
	strainMiner	**50.1**	56	335	**81.65**	72.73
E. coli	metaFlye	6.05	374	22	216.18	216.53
2 strains	Strainberry	10.8	**1152**	11	24.83	64.73
(9.9 Mb)	HairSplitter	**10.4**	230	47	53.75	62.67
	strainMiner	**9.65**	45	**8**	**15.4**	**59.77**

A limitation of the strainMiner pipeline, for now, is the limited contiguity of the assembly produced. Adapting the pipeline to the specificities of the strainMiner module would be the first lead to pursue to improve this. Indeed, the assemblies obtained through the strainMiner pipeline had often much lower contiguity compared to the one obtained by HairSplitter, suggesting that the properties of the read separation computed by the strainMiner module and the native HairSplitter read-separation module may be quite different. Tailoring a scaffolding step could thus improve contiguity.

One aspect of strainMiner that would require

further investigation is its approach of dividing the problem into non-overlapping fixed-length windows. While this division results in a clean input matrix for the Integer Linear Programming (ILP) solver, it introduces an arbitrary separation of loci that may jointly carry valuable information. This division diminishes the sensitivity of strainMiner when dealing with highly similar strains, as the number of loci in a single window may be too low for detection, while there may be enough loci overall to detect the strain. A possible solution to address this limitation would be considering overlapping windows. Another option would be the adaptation of the model to handle miss-

ing values without the need for imputation, which for now limits the length of the windows.

Finally, users might find themselves limited by the Gurobi license. Academic license is free but limited to three instances of Gurobi running at the same time, limiting the multi-threading potential. Attempts to use the free CBC solver showed a decrease in performance.

SOFTWARE AVAILABILITY

strainMiner is freely availabe on github at github.com/RolandFaure/strainMiner, with the Affero GPL3 license.

ACKNOWLEDGEMENTS

We wish to thank Dominique Lavenier, who formulated the first version of the optimization problem.

Many thanks to Riccardo Vicedomini for his help in pre-reviewing the article.

ChatGPT was used to correct and reformulate the writing of the article.

REFERENCES

Albanese, D. and Donati, C. (2017). Strain profiling and epidemiology of bacterial species from metagenomic sequencing. *Nature Communications*, 8.

Baaijens, J., Aabidine, A., Rivals, E., and Schönhuth, A. (2017). De novo assembly of viral quasispecies using overlap graphs. *Genome Research*.

Bansal, V. (2022). Hapcut2: A method for phasing genomes using experimental sequence data. *Methods in molecular biology*, 2590:139–147.

Bertrand, D., Shaw, J., Kalathiyappan, M., Ng, A. H. Q., Kumar, M. S., Li, C., Dvornicic, M., Soldo, J. P., Koh, J. Y., Tong, C., Ng, O. T., Barkham, T., Young, B., Marimuthu, K., Chng, K. R., Sikic, M., and Nagarajan, N. (2019). Hybrid metagenomic assembly enables high-resolution analysis of resistance determinants and mobile elements in human microbiomes. *Nature Biotechnology*, 37(8):937–944.

Bickhart, D., Kolmogorov, M., Tseng, E., Portik, D., Korobeynikov, A., Tolstoganov, I., Uritskiy, G., Liachko, I., Sullivan, S., Shin, S., Zorea, A., Pascal, V., Panke-Buisse, K., Medema, M., Mizrahi, I., Pevzner, P., and Smith, T. (2022). Generating lineage-resolved, complete metagenome-assembled genomes from complex microbial communities. *Nature Biotechnology*, 40.

Cleary, B., Brito, I., Huang, K., Gevers, D., Shea, T., Young, S., and Alm, E. (2015). Detection of low-abundance bacterial strains in metagenomic datasets

by eigengenome partitioning. *Nature biotechnology*, 33.

Faure, R., Flot, J.-F., and Lavenier, D. (2023). Hairsplitter: separating strains in metagenome assemblies with long reads. In *Proceedings of JOBIM 2023*, pages 124–131.

Fedarko, M., Kolmogorov, M., and Pevzner, P. (2022). Analyzing rare mutations in metagenomes assembled using long and accurate reads. *Genome research*, 32.

Feng, X., Cheng, H., Portik, D., and Li, H. (2022). Metagenome assembly of high-fidelity long reads with hifiasm-meta. *Nature Methods*, 19:1–4.

Feng, Z., Clemente, J., Wong, B., and Schadt, E. (2021). Detecting and phasing minor single-nucleotide variants from long-read sequencing data. *Nature Communications*, 12:3032.

Fix, E. and Hodges, J. L. (1989). Discriminatory analysis. nonparametric discrimination: Consistency properties. *International Statistical Review / Revue Internationale de Statistique*, 57(3):238–247.

Frank, C., Werber, D., Cramer, J. P., Askar, M., Faber, M., an der Heiden, M., Bernard, H., Fruth, A., Prager, R., Spode, A., Wadl, M., Zoufaly, A., Jordan, S., Kemper, M. J., Follin, P., Müller, L., King, L. A., Rosner, B., Buchholz, U., Stark, K., and Krause, G. (2011). Epidemic profile of shiga-toxin–producing escherichia coli o104:h4 outbreak in germany. *New England Journal of Medicine*, 365(19):1771–1780. PMID: 21696328.

Gurobi Optimization, LLC (2023). Gurobi [computer software]. gurobi.com.

Kang, X., Luo, X., and Schönhuth, A. (2022). StrainXpress: strain aware metagenome assembly from short reads. *Nucleic Acids Research*, 50(17):e101–e101.

Kazantseva, E., Donmez, A., Pop, M., and Kolmogorov, M. (2023). stRainy: assembly-based metagenomic strain phasing using long reads. preprint, Bioinformatics.

Kolmogorov, M., Bickhart, D. M., Behsaz, B., Gurevich, A., Rayko, M., Shin, S. B., Kuhn, K., Yuan, J., Polevikov, E., Smith, T. P. L., and Pevzner, P. A. (2020). metaFlye: scalable long-read metagenome assembly using repeat graphs. *Nature Methods*, 17(11):1103–1110.

Koren, S., Walenz, B. P., Berlin, K., Miller, J. R., Bergman, N. H., and Phillippy, A. M. (2017). Canu: scalable and accurate long-read assembly via adaptive k-mer weighting and repeat separation. *Genome Research*, 27(5):722–736.

Mapleson, D., Accinelli, G., Kettleborough, G., Wright, J., and Clavijo, B. (2016). Kat: A k-mer analysis toolkit to quality control ngs datasets and genome assemblies. *Bioinformatics (Oxford, England)*, 33.

Mikheenko, A., Saveliev, V., and Gurevich, A. (2015). Metaquast: Evaluation of metagenome assemblies. *Bioinformatics*, 32:btv697.

Peeters, R. (2003). The maximum edge biclique problem is NP-complete. 131:651–654.

Quince, C., Delmont, T. O., Raguideau, S., Alneberg, J., Darling, A. E., Collins, G., and Eren, A. M. (2017).

DESMAN: a new tool for de novo extraction of strains from metagenomes. *Genome Biology*, 18(1):181.

Quince, C., Nurk, S., Raguideau, S., James, R., Soyer, O. S., Summers, J. K., Limasset, A., Eren, A. M., Chikhi, R., and Darling, A. E. (2020). Metagenomics Strain Resolution on Assembly Graphs. preprint, Bioinformatics.

Rodriguez Jimenez, A., Guiglielmoni, N., Goetghebuer, L., Dechamps, E., George, I., and Flot, J.-F. (2022). Comparative genome analysis of vagococcus fluvialis reveals abundance of mobile genetic elements in sponge-isolated strains. *BMC Genomics*, 23.

Troyanskaya, O., Cantor, M., Sherlock, G., Brown, P., Hastie, T., Tibshirani, R., Botstein, D., and Altman, R. B. (2001). Missing value estimation methods for dna microarrays. *Bioinformatics*, 17(6):520–525.

Truong, D. T., Tett, A., Pasolli, E., Huttenhower, C., and Segata, N. (2017). Microbial strain-level population structure and genetic diversity from metagenomes. *Genome Research*, 27:gr.216242.116.

Vicedomini, R., Quince, C., Darling, A. E., and Chikhi, R. (2021). Strainberry: automated strain separation in low-complexity metagenomes using long reads. *Nature Communications*, 12(1):4485.

Wick, R. (2019). Badread: simulation of error-prone long reads. *Journal of Open Source Software*, 4(36):1316.